

The Stratego Tutorial

Eelco Visser

STRATEGO

For Stratego version 0.4.18

www.stratego-language.org

Department of Information and Computing Sciences
Universiteit Utrecht
September 29, 2000

Copyright © 1998, 1999, 2000 Eelco Visser

Permission is granted to make and distribute verbatim copies of this manual provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this manual under the conditions for verbatim copying, provided that they are marked clearly as modified versions, that the author's names and title are unchanged (though subtitles and additional authors' names may be added), and that other clearly marked sections held under separate copyright are reproduced under the conditions given within them, and that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this manual into another language, under the above conditions for modified versions, except that this permission notice may be stated in a translation approved by the Free Software Foundation.

Address:
Department of Computer Science
Universiteit Utrecht
P.O.Box 80089
3508 TB Utrecht
email: visser@acm.org
<http://www.cs.uu.nl/~visser/>

The Stratego Tutorial

Eelco Visser

September 29, 2000

Summary

Stratego is a language for the specification of transformation rules and strategies for applying them. Specifications consist of a collection of modules that define the signature of the object language(s) of the transformation, transformation rules and strategies. The Stratego compiler translates specifications to C code. Together with a provided run-time system these generated programs can be used to apply the specified transformations.

This documents gives a tutorial for Stratego. It introduces the constructs of the language, presents common architectures for transformation systems, outlines the library of specifications that comes with the language, explains several example specifications and a instructions for installation and usage of the implementation.

CONTENTS

STRATEGO AT A GLANCE	2
SPECIFYING PROGRAM TRANSFORMATION SYSTEMS WITH STRATEGO	2
WHEN TO USE AND NOT TO USE STRATEGO	4
STRATEGO FEATURES	6
USING THE STRATEGO DOCUMENTATION	7
INSTALLATION	8
DOWNLOADING STRATEGO	8
INSTALLING STRATEGO	10
STRUCTURE OF THE DISTRIBUTION	12
BOOTSTRAPPING THE STRATEGO COMPILER	14
STRATEGO BY EXAMPLE	16
DEVELOPING TRANSFORMATION SYSTEMS WITH STRATEGO	16
Writing Specifications	18
DESCRIBING TERMS WITH SIGNATURES	18
SPECIFYING BASIC TRANSFORMATIONS WITH RULES	20
SIMPLIFICATION RULES FOR PROPOSITIONAL FORMULAE	22
COMBINING RULES INTO TRANSFORMATION SYSTEMS WITH STRATEGOES	24
Compiling Specifications	26
COMPILING SPECIFICATIONS TO EXECUTABLE PROGRAMS	26
disjnf	27
DETECTING ERRORS IN STRATEGO SPECIFICATIONS	28
DEBUGGING STRATEGO PROGRAMS	29
FURTHER READING	30
EXTENDING YOUR UNDERSTANDING OF STRATEGO	30
BIBLIOGRAPHY	31

SPECIFYING PROGRAM TRANSFORMATION SYSTEMS WITH STRATEGO

Stratego is a language for the specification of program transformation systems based on the paradigm of rewriting strategies. Stratego specifications give rise to term transformers that read a term and output a transformed term. Such transformation components can easily be composed into transformation pipelines.

TERM TRANSFORMATION SYSTEMS

Stratego is a language for the specification of automatic program transformation systems. In the Stratego model, a program transformation system is a program that reads a program, applies a transformation to it, and, if that succeeds, spits out the transformed program or otherwise states that it failed (Figure 1). Programs are supposed to be represented by means of abstract syntax trees or *terms*. A transformation system is thus a program that transforms terms to terms.

Transformation components can be composed into pipelines that interface by means of intermediate terms (Figure 2). In this fashion large program transformation systems can be built from small reusable components, provided that the components synchronize on the intermediate languages used.

SPECIFYING TRANSFORMATIONS WITH REWRITING STRATEGIES


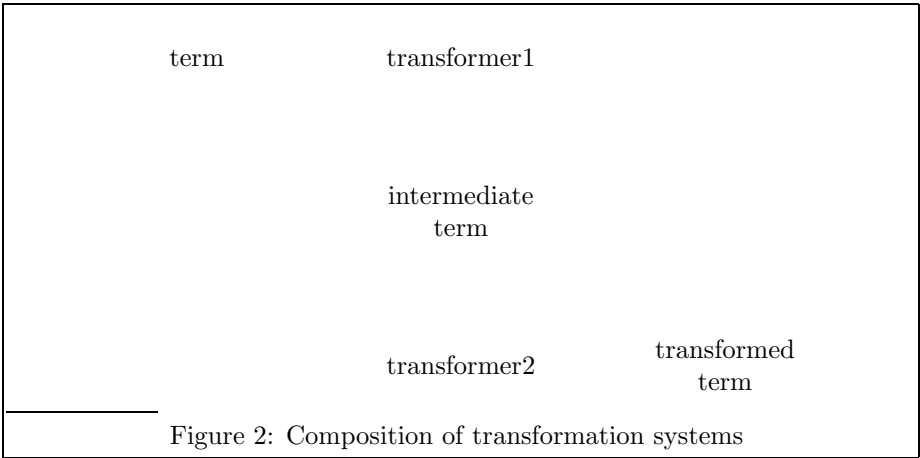

Basic transformation rules can be expressed as *term rewrite rules* of the form $l \rightarrow r$ that transform a term matching l to an instantiation of r .

Pure rewrite rules with a standard evaluation strategy are not sufficient for most program transformation jobs, because more fine grained control over their application is needed. Usually such control is encoded in the rewrite rules themselves, leading to bad reusability of rules and a large overhead for expressing traversals over the abstract syntax tree.

Instead Stratego offers a language of user-definable *rewriting strategies*¹ that provide control over the application of transformation rules. This is achieved by a set of operators for the expression of generic traversals over abstract syntax trees (Figure 3). The basic operators can be combined to achieve a wide range of traversal algorithms.

The use of strategies greatly reduces the overhead caused by explicitly programming term traversals. Furthermore, since individual transformation rules are freed from the burden of encoding traversals, they can be reused in other situations. Finally, the strategy operators turn out to be an interesting new programming paradigm that allow one to reach deep into terms without much effort.

¹The name Stratego is derived from this feature of the language.



WHEN TO USE AND NOT TO USE STRATEGO

Stratego is designed for the scheduling of the application of a series of local transformation rules on terms. Therefore, the language is primarily useful in source to source program transformation, application generation and code generation. It is not designed for theorem proving, program analysis, interactive program derivation or general purpose programming, although it can be used to assist in these activities.

WHAT STRATEGO IS FOR

Stratego is designed for the specification of automatic program transformation systems. Typical application areas of automatic program transformation are source-to-source transformations, application generation, code generation, program derivation and documentation generation.

Source-to-source transformation is the transformation of a program to a program in the same language. Examples of source to source transformations are simplification, optimization and desugaring, i.e., the elimination of syntactic abstractions.

Application generation is the translation of a program in a high-level (end-user) application language to an implementation in a general purpose programming language.

Code generation is the translation of program constructs to sequences of machine code instructions.

Automatic program derivation is concerned with the generation of programs from data. Examples of program derivations are the derivation of a constructor signature from a context-free grammar and the derivation of a pretty-printer from a context-free grammar [1].

Documentation generation is the extraction of documentation information from a program.

Program analysis is concerned with deriving properties about elements of programs. Examples of program analysis are type checking, type inference, and data flow analysis. Although Stratego can be used to implement such analyses, it is not particularly well supported. This aspect should probably be strengthened.

WHAT STRATEGO IS NOT FOR

Stratego is not necessarily applicable in other areas of program manipulation related applications. In particular, it is probably not the right platform for interactive program transformation, theorem proving and general purpose programming.

Interactive program transformation is concerned with the gradual transformation from a high-level specification to an implementation by intervention from a programmer. The programmer has to decide what refinement steps to apply to what part of the specification. However, Stratego can be used to make program

transformation scripts that automatically transform a specific specification to an implementation. Such scripts can be written incrementally to reflect the increased understanding of the refinement.

Theorem proving is used for establishing correctness of programs. Stratego does not support theorem proving. However, it can be used to generate proof obligations as input for theorem provers.

Stratego is not a general purpose programming language to use for application programming, user-interface implementation, etc. However, it can very well be used to generate or transform programs in a general purpose language.

- Abstract and concrete syntax tree manipulation
- Simplification
- Optimization
- Desugaring (eliminating syntactic abstractions)
- Code generation
- Application generation
- Documentation generation
- Program analysis

Figure 4: Application areas

- Theorem proving
- Interactive program transformation / derivation
- General purpose programming

Figure 5: Non application areas

STRATEGO FEATURES

- Modular specification of transformation systems
- Separation of rules and strategies
- Reuse of rules
- Primitive operators of rewriting
- Concise specification of term traversals
- Administration of failure
- Generic operations on (anonymous) trees
- Persistent terms

Figure 6: Key features of Stratego

- Higher-order abstract syntax

Figure 7: Features not supported by Stratego

USING THE STRATEGO DOCUMENTATION

To support the production of Stratego specifications, the distribution comes with four documents: a tutorial, an overview of the library, a reference manual and the documentation of the compiler.

THE STRATEGO TUTORIAL

The Stratego Tutorial shows how to install and use the Stratego compiler and gives basic examples of Stratego specifications. The tutorial is basic reading material for every Stratego user.

THE STRATEGO LIBRARY

The Stratego Library documents all generic rules and strategies in the standard library that comes with the distribution. The Library includes the declaration of all primitive operations (I/O, string manipulation, arithmetic, process creation, etc.). Refer to the library to find common strategies and to understand the primitive operations.

THE STRATEGO REFERENCE MANUAL

The Stratego Reference manual documents all constructs of the language by giving their syntax and a description of the semantics. Refer to the reference manual to clarify syntactic issues and understand the concepts of the semantics.

THE STRATEGO COMPILER

The Stratego Compiler contains the complete annotated source of the compiler and run-time system. The compiler should be read by Stratego developers and those that want to get a detailed understanding of the implementation of the operational semantics. It should not be necessary to read the Compiler for ordinary use. However, it is a case study of the use of Stratego and in which a number of idioms are used.

OTHER STRATEGO PUBLICATIONS

Several other publications introduce Stratego, discuss applications, semantics and related work. See for example [?, 2, 3, 7, 8].

DOWNLOADING STRATEGO

Stratego is available on the Internet. The Stratego distribution contains the sources of the Stratego compiler and the Stratego compiler. Utilities for use with Stratego are available in separate packages.

THE STRATEGO PACKAGE

The Stratego compiler is implemented in Stratego itself. The distribution contains the bootstrapped C sources of the compiler, the implementation of the run-time system, the Stratego library, the Stratego sources of the compiler, and the Stratego documentation.

DOWNLOADING

The latest release of the Stratego compiler can be found at the Stratego web page: <http://www.cs.uu.nl/~visser/stratego/>. Download the latest release `stratego-x.x.tar.gz`

REQUIREMENTS

The following packages are required for the installation of Stratego:

- ATerm library
- gmake
- gcc [developed with version egcs-2.91.66]
- bison, flex

Except for the ATerm library, these are all standard elements of a GNU installation.

Starting with version 0.4.4a, the ATerm library is no longer bundled with the Stratego distribution. It is available via the Stratego web page though.

The distribution for Stratego version 0.4.18 has been tested on Linux (RedHat 6.0) on a Pentium II. Previous versions have been tested also on Sun4 and Solaris. Installation under IRIX at SGI machines has been problematic for previous versions. Installation under Windows NT with Cygnus has not yet been attempted. I would be interested in hearing any reports from such an attempt.




Figure 8: The Stratego webpage [<http://www.cs.uu.nl/~visser/stratego/>] contains the distribution of the compiler, publications about Stratego.

INSTALLING STRATEGO

The Stratego distribution has been prepared with autoconf and automake and should be portable to any platform with the required packages. Installation requires only the configuration of the package using a configuration script.

INSTALLING THE ATERM LIBRARY

Choose a location for the installation of the aterm package, say `/loc/of/aterm/`.

```
> cd /loc/of/aterm/
```

Download the ATerm package `aterm-x.y.tar.gz`. Unpack it using the command

```
> tar xzf aterm-x.y.tar.gz
```

Configure the aterm package, setting the prefix to the directory where the library should be installed, i.e., if prefix points to directory `dir`, then the library will be installed in `dir/lib` and the executables in `dir/bin`.

```
> cd aterm-x.y
```

```
> ./configure --prefix=/installation/of/aterm
```

make

```
> gmake
```

and install

```
> gmake install
```

UNPACKING THE STRATEGO PACKAGE

Choose a location for the installation of Stratego, say `/loc/of/stratego/`.

```
> cd /loc/of/stratego/
```

Download the Stratego package `stratego-x.y.z.tar.gz`. Unpack it using the command:

```
> tar xzf stratego-x.y.z.tar.gz
```

```
> cd stratego-x.y.z/
```

CONFIGURING THE STRATEGO PACKAGE

If you only intend to use the Stratego compiler and not work on its development set prefix to any appropriate value when installing, e.g.,

```
> ./configure --prefix=/usr/local \
              --with-aterm=/installation/of/aterm
```

The directory `/installation/of/aterm` should be the same as the prefix you configured the aterm library with. If prefix points to directory `dir`, then the Stratego compiler will be installed in `dir/bin`, the libraries in `dir/lib` and the Stratego library in `dir/share/stratego`.

If you intend to change the compiler specification in `spec/` and bootstrap the compiler, you need to set prefix to the local directory. Or at least a directory where you have write permissions such that it is easy to re-install the compiler. This is the what I usually do:

```
> ./configure --prefix='pwd' \
              --with-aterm=/installation/of/aterm
```

This entails that you need to include `'pwd'/bin` in your `PATH`.

Chapter: Installation

Note that this last option may also be needed if you don't have permissions to install in standard installation directories.

MAKING AND INSTALLING

For both kinds of configurations, make the compiler by

```
> gmake
> gmake install
```

SETTING THE PATH

Finally, extend the definition of your path to point to the stratego compiler. In a user configuration, set

```
> PATH=/usr/local/bin:$PATH
```

In a developer configuration, set

```
> PATH=/loc/of/stratego/bin:$PATH
```

CHECKING THE INSTALLATION

Now check that the compiler is installed by typing

```
> which sc
```

which should return something like

```
/loc/of/stratego/bin/sc
```

```
> cd /loc/of/stratego/
> tar zxf stratego-x.y.z.tar.gz
> cd stratego-x.y.z
> ./configure --prefix=/usr/local \
               --with-aterm=/installation/of/aterm
> gmake
> gmake install
```

Figure 9: Summary of the installation procedure

Chapter: Installation

STRUCTURE OF THE DISTRIBUTION

(*** TODO: OVERVIEW OF THE DISTRIBUTION + RATIONALE ***)

- spec : sources of the compiler components
 - back : back-end
 - form : format checking
 - front : front-end
 - lib : library of the compiler
 - match : matching automaton
 - opt : optimizer
 - post : post-processing
 - pp : pretty-printing
 - rts : run-time system
 - sc : the Stratego compiler; glues components together
 - sig : signatures
 - slib : Stratego library
 - * spec : specifications
 - * src : C implementation of primitives
 - syn : syntax
 - test : regression tests
- src : the bootstrapped compiler components
 - a mirror of the spec tree with only the generated C code, without specifications
- doc : documentation
 - compiler
 - library
 - reference
 - tutorial
- share
 - tex : LaTeX packages
 - stratego : installed Stratego library
 - bib : BibTeX files
- www : Stratego web page
- bin : installed executables
- include : installed include files
- lib : installed binary libraries

Figure 10: Structure of the Stratego distribution

BOOTSTRAPPING THE STRATEGO COMPILER

The Stratego compiler is defined in Stratego. The distribution contains the complete Stratego source of the compiler. This makes it possible to adapt the language and/or its implementation. In order to obtain a new working compiler it is necessary to follow a careful bootstrapping procedure.

To create a new version of the compiler, edit the appropriate files in the `spec/*` directories to reflect the change in the run-time system, compiler, and/or parser. Then do the following

```
> cd spec
> make
> make install
```

Then use the generated compiler to compile some test specifications. If that seems to work, try to bootstrap by

```
> make bootstrap
```

This will compile the compiler components with itself.

If this goes wrong at some point, you'll have to reinstall the stable compiler in the `src/` tree:

```
> cd ..
> make install
```

It can be useful to have two installations of the compiler, one of them the stable compiler from `src`. To achieve this, do the following.

- Make a user configuration as described above and make and install the compiler with this configuration. This puts a stable version of the compiler in a standard place like `/usr/local/bin`.
- Make a developer configuration as described above.
- In one terminal use the stable compiler to compile the compiler components, i.e., `PATH=/usr/local/bin:$PATH`
- In another terminal use the newly compiled compiler to test your changes i.e., `PATH=/path/to/stratego/bin:$PATH`

Now if something goes wrong it is easier to recompile the compiler components.

(*** TODO: PICTURE OF EFFECTS OF VARIOUS MAKE TARGETS
***)

Figure 11: The bootstrapping procedure

DEVELOPING TRANSFORMATION SYSTEMS WITH STRATEGO

A Stratego specification consists of a collection of modules that contain signatures, rules and strategies that together define a transformation system. The Stratego compiler translates a specification to an executable program that can be used to transform terms.

A Stratego specification defines a transformation on terms. A specification consists of a signature describing the structure of the language, rules that define the steps of the transformation, and strategies that combine the rules in a transformation system.

All these elements can be mixed in any order in one big file. However, to make specifications reusable, Stratego provides a module system that can be used to spread definitions over several files. Thus, a specification consists of a collection of modules that define signatures, rules and strategies.

It is usually a good idea to divide a specification over several modules: one defining the language, several defining various sorts of transformation rules and strategies, and one defining the main strategy that applies these rules.

An executable transformation system can be derived from a specification by means of the Stratego compiler *sc*. Compiled specifications can be used to transform terms.

Specifications are rarely correct in one go. Some errors are caught by the compiler, others have to be detected by inspection of the specification and by debugging techniques.

EXAMPLE: SIMPLIFICATION OF PROPOSITIONAL FORMULAE

As an example, consider the simplification of propositional formulae. Figure 13 shows the import graph for the specification. An arrow points from the importing module to the module that is imported. Module **prop** defines the signature of propositional formulae. Module **prop-truth** defines the truth rules for the propositional connectives. Module **prop-laws** defines a number of rules defining sound transformations on formulae. Module **prop-simplify** imports modules **prop-laws** and **prop-truth** and defines various ways of combining the rules in simplifying transformations. Module **disjnf** declares one of these transformations, i.e., the simplification of a formula to disjunctive normal form, to be the main transformation.

In addition, module **prop-simplify** imports the module **simple-traversal**, which defines a number of standard term traversals. Module **disjnf** imports **io**, which defines strategies for term input and output. The latter two modules are part of the Stratego library.

- Develop specification = collection of modules defining
 - Signature
 - Rules
 - Strategies
- Compile specification
 - Repair syntactic errors
- Use compiled specification
 - Detect logical errors

Figure 12: The development process




Figure 13: Import graph of a specification

DESCRIBING TERMS WITH SIGNATURES

Stratego specifications define transformations on terms. The structure of terms can be described by means of signatures.

TERMS

In the Stratego model, programs and all other objects that are transformed are represented by means of terms. In their most basic form, terms are described by the rule $t := C(t_1, \dots, t_n)$. That is, a term is an application of a constructor C to zero or more other terms.

An example term is `Not(Or(And(Atom("B"), Atom("A")), Atom("A")))`

TREE REPRESENTATION OF TERMS

A term can be represented by means of a tree diagram. Figure 16 shows a tree diagram for the example term above. Tree diagrams are a model for the representation of terms in a computer's memory. Nodes correspond to tree cells that contain a node tag and pointers (arrows) to the cells corresponding to the sub-terms of the node.

DAG REPRESENTATION OF TERMS

In fact, in the Stratego implementation terms are represented by means of directed, acyclic graphs (dags) that share sub-terms. That is, occurrences of the same sub-term are actually represented by the same cell in memory. For instance, the diagram in Figure 17 shows the sharing of the sub-term `Atom("A")` by the `And` and the `Or` node.

SIGNATURES

A signature is used to describe the names of term constructors and the number and type of their arguments. An operator declaration $C : s$ defines a nullary constructor C . An operator declaration $C : s_1 * \dots * s_n \rightarrow s$ defines an n -ary constructor C . A signature characterizes a sub-set of the universal set of terms described above.

For example, Figure 15 shows module `prop` (in file `prop.r`), which defines the structure of propositional formulae. The signature defines the sort `Prop` of propositional formulae. The operations or constructors of the language are `Atom` that constructs propositional letters, `Not` for negation, `And` for conjunction, `Or` for disjunction, `Impl` for implication and `Eq` for equivalence. The term in Figure 14 is an example of a formula over this signature.

```
Not(And(Not(Or(Atom("A"), Not(And(Atom("B"), Atom("C"))))),
      Atom("D"))))
```

Figure 14: A term representing a propositional formula (file: `prop.trm`).

```

module prop
signature
  sorts Prop
  constructors
    False : Prop
    True  : Prop
    Atom  : String -> Prop
    Not   : Prop -> Prop
    And   : Prop * Prop -> Prop
    Or    : Prop * Prop -> Prop
    Impl  : Prop * Prop -> Prop
    Eq    : Prop * Prop -> Prop

```

Figure 15: Signature of propositional formulae (file: prop.r).




Figure 16: Tree structure of terms




Figure 17: Directed-acyclic graph structure of terms

SPECIFYING BASIC TRANSFORMATIONS WITH RULES

Transformation rules are the basic components of transformations defined in Stratego. A rule transforms a term at the root if the term matches the left-hand side of the rule. The result is the instantiation of the right-hand side of the rule.

TRANSFORMATION

A transformation is a modification of a term. An example of a transformation on propositional formulae is the transformation of an arbitrary formula to disjunctive normal form.

RULE

Transformations can be achieved by consecutively applying a number of small modifications to a term. In Stratego, a transformation rule defines a single transformation step. A rule consists of a label (e.g., `DMO`), a left-hand side term pattern (e.g., `Not(Or(x, y))`) and a right-hand side term pattern (e.g., `And(Not(x), Not(y))`).

Notation: `DMO : Not(Or(x, y)) -> And(Not(x), Not(y))`

TERM PATTERNS ARE TERMS WITH VARIABLES

A term pattern is a term with variables. That is, term patterns are described by the rule `tp := x | C(tp1, ..., tpn)`, where `x` is a variable. Variables do not have to be declared. All nullary (i.e., non-applied) constructors that are not declared in the signature are assumed to be variables.

APPLYING TRANSFORMATION RULES

A rule defines a one-step transformation *at the root* of a term. A rule applies to a term if the term matches the left-hand side pattern of the rule. If that is the case the term is replaced with the right-hand side of the term in which the variables are replaced by the corresponding sub-terms that were matched in the left-hand side. Thus a transformation rule succeeds to apply if the left-hand side matches, and fails to apply if that is not the case. In general, a transformation may succeed or fail to apply to a term.

For example, the De Morgan rule `DMO` transforms the term

```
Not(Or(And(Atom("B"), Atom("A")), Atom("A")))
```

to

```
And(Not(And(Atom("B"), Atom("A"))), Not(Atom("A")))
```

Figure 18 illustrates the effect of this transformation on the tree structure of a term.

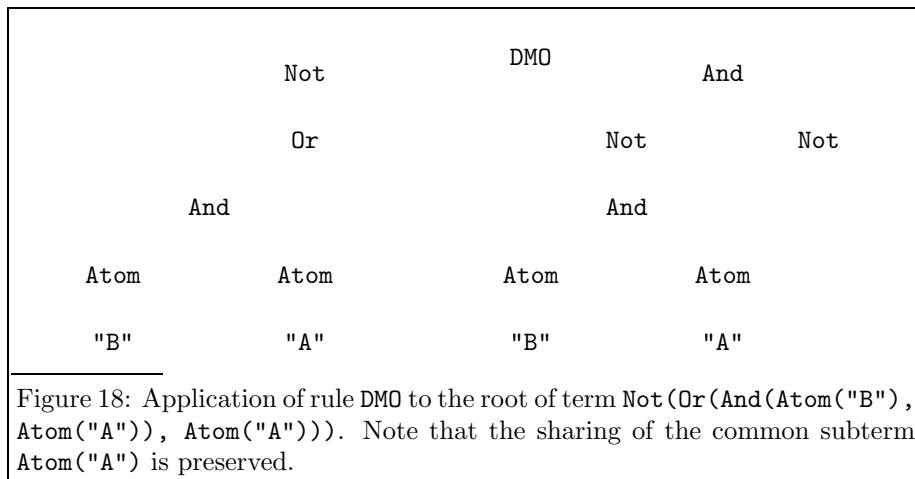


Figure 18: Application of rule DMO to the root of term `Not(Or(And(Atom("B")), Atom("A")), Atom("A"))`. Note that the sharing of the common subterm `Atom("A")` is preserved.

SIMPLIFICATION RULES FOR PROPOSITIONAL FORMULAE

There are a great number of valid transformation rules on propositional formulae. Although each is useful in itself, together they form a non-terminating rewrite system.

Propositional formulae can be simplified using well-known rules such as De Morgan's laws and distribution laws. Some of these simplification rules are defined in module `prop-laws` (Figure 20). Module `prop-truth` (Figure 19) defines rules for constant propagation in formulae.

Note that these rules give rise to a number of possible normalizations of formulae. Interpretation of the rules as a rewrite system is not useful since they define a non-terminating system. For example, rules DAOL and DOAL can be applied alternately without reaching a normal-form.

```

module prop-truth
imports prop

rules

  T1  : Not(True) -> False      T2  : Not(False) -> True

  T3  : And(True, x) -> x      T5  : And(False, x) -> False
  T4  : And(x, True) -> x      T6  : And(x, False) -> False

  T7  : Or(True, x) -> True    T9  : Or(False, x) -> x
  T8  : Or(x, True) -> True    T10 : Or(x, False) -> x

  T11 : Impl(True, x) -> x     T14 : Impl(False, False) -> True
  T12 : Impl(x, True) -> True  T15 : Impl(True, False) -> False
  T13 : Impl(False, x) -> True

  T16 : Eq(False, x) -> Not(x) T18 : Eq(True, x) -> x
  T17 : Eq(x, False) -> Not(x) T19 : Eq(x, True) -> x

```

Figure 19: Truth rules for propositional formulae.

```

module prop-laws
imports prop
rules

AA    : And(And(x, y), z) -> And(x, And(y, z))
AO    : Or(Or(x, y), z) -> Or(x, Or(y, z))
AI    : Impl(Impl(x, y), z) -> Impl(x, Impl(y, z))
AE    : Eq(Eq(x, y), z) -> Eq(x, Eq(y, z))

CA    : And(x, y) -> And(y, x)
CO    : Or(x, y) -> Or(y, x)
CE    : Eq(x, y) -> Eq(y, x)

IDA   : And(x, x) -> x
IDO   : Or(x, x) -> x
IDI   : Impl(x, x) -> True
IDE   : Eq(x, x) -> True

DAOL  : And(Or(x, y), z) -> Or(And(x, z), And(y, z))
DAOR  : And(z, Or(x, y)) -> Or(And(z, x), And(z, y))
DOAL  : Or(And(x, y), z) -> And(Or(x, z), Or(y, z))
DOAR  : Or(z, And(x, y)) -> And(Or(z, x), Or(z, y))

DN    : Not(Not(x)) -> x

DMA   : Not(And(x, y)) -> Or(Not(x), Not(y))
DMO   : Not(Or(x, y)) -> And(Not(x), Not(y))
DMI   : Not(Impl(x, y)) -> And(x, Not(y))
DME   : Not(Eq(x, y)) -> Or(And(Not(x), y), And(x, Not(y)))

DefN  : Not(x) -> Impl(x, False)
DefI  : Impl(x, y) -> Or(Not(x), y)
DefE  : Eq(x, y) -> And(Impl(x, y), Impl(y, x))
DefO1 : Or(x, y) -> Impl(Not(x), y)
DefO2 : Or(x, y) -> Not(And(Not(x), Not(y)))
DefA1 : And(x, y) -> Not(Or(Not(x), Not(y)))
DefA2 : And(x, y) -> Not(Impl(x, Not(y)))

IDefI : Or(Not(x), y) -> Impl(x, y)
IDefE : And(Impl(x, y), Impl(y, x)) -> Eq(x, y)

```

Figure 20: Simplification rules for propositional formulae.

COMBINING RULES INTO TRANSFORMATION SYSTEMS WITH STRATEGIES

Applying all transformation rules until none applies anymore, is often not a good strategy in program transformation because the transformation might not terminate or because many different transformation paths are possible, one of which needs to be chosen. Programmable strategies provide a way to control the application of transformation rules.

Standard term rewriting apply rules using a default strategy in which all rules are applied as long as possible. Therefore, for a rewrite system to be meaningful it should be terminating and confluent, i.e., whatever transformation path is chosen the same term results.

In general, however, given a collection of meaningful transformation rules many possible transformations on a term are possible. Depending on the desired outcome, different rules should be applied.

For example, the rules for transformation of propositional formulae can be used to transform formulae to a wide variety of normal forms. Formulae in *disjunctive normal form* are disjunctions of conjunctions of atoms or negations of atoms. Furthermore, disjunctive normal forms do not contain implications, equivalences or truth values. A formula in *conjunctive normal form* is a conjunction of disjunctions of atoms or negations of atoms. Finally, any propositional formula can be expressed using just implication and **False**.

Each of these normal forms can be achieved by combining a different selection of the rules from module **prop-laws** into a transformation.

TRANSFORMATION STRATEGIES

Stratego supports user-definable strategies for the application of transformation rules. A strategy is a program that specifies which rules to apply in what order to which sub-terms of a term.

Rule labels are basic strategies. Strategies are composed from rule labels with a number of strategy operators. For example, the specification in Figure 22 uses non-deterministic choice (+), sequential composition (;), and recursive closure (**rec** *x*(*s*)). Application of transformations below the root of a term is achieved by means of traversal operators such as **all**(*s*), which applies the strategy *s* to all direct sub-terms of the term to which it is applied.

In module **prop-simplify** the rules from modules **prop-laws** and **prop-truth** are combined into strategies that simplify formulae. Strategy **disjnf** transforms propositional formulae to disjunctive normal form. It defines a traversal over a term that pushes **Not**s inwards on the way down and pushes **And**s inwards on the way up. Strategy **conjnf** transforms formulae to conjunctive normal form in a similar way, but using **Or**-distributivity instead of **And**-distributivity. Strategy **desugar** defines implication and equivalence in terms of conjunction, disjunction, and negation.

```

module simple-traversal
strategies

  try(s)          = s <+ id
  repeat(s)       = rec x(try(s; x))

  topdown(s)      = rec x(s; all(x))
  bottomup(s)     = rec x(all(x); s)
  downup(s)       = rec x(s; all(x); s)
  downup2(s1, s2) = rec x(s1; all(x); s2)

  oncetd(s)       = rec x(s <+ one(x))
  oncebu(s)       = rec x(one(x) <+ s)

  alltd(s)        = rec x(s <+ all(x))

  (* etc. *)

```

Figure 21: Simple traversal strategies. This is an extract from the library module `simple-traversal`. See the Stratego Library for more example strategies.

```

module prop-simplify
imports prop-laws prop-truth simple-traversal
strategies

  T          = T1 + T2 + T3 + T4 + T5 + T6 + T7 + T8 + T9 + T10 +
              T11 + T12 + T13 + T14 + T15 + T16 + T17 + T18 + T19

  desugar    = topdown(try(DefI + DefE))
  eval       = bottomup(repeat(T))
  impl-nf    = topdown(repeat(DefN + DefA2 + DefO1 + DefE))

  disj-nf    = innermost(DAOL + DAOR + DN + DMA + DMO)
  conj-nf    = innermost(DOAL + DOAR + DN + DMA + DMO)

```

Figure 22: Simplification strategies for propositional formulae.

COMPILING SPECIFICATIONS TO EXECUTABLE PROGRAMS

A specification defines a transformation. To make the transformation executable on a computer, it is necessary to compile it to a standard programming language. The Stratego compiler translates a Stratego specification to a stand-alone C program.

THE MAIN STRATEGY

In order to compile a specification it is necessary to indicate the main transformation strategy that the compiled program should apply to terms. The compiler assumes that a definition for the strategy `main` exists that indicates the transformation to perform. For example, module `disjnf` in Figure 23 defines the strategy `main` to first desugar the formula and then transform it to disjunctive normal form.

Note the use of the strategy `stdio` in the definition of `main`. Input/output behaviour needs to be defined in the specification itself, i.e., it is not handled in a default manner. The `stdio` strategy implements the most basic interface, i.e., read a term from standard input and write the result to standard output. `stdio` is defined in the library module `io`.

COMPILING SPECIFICATIONS

A specification is compiled with the Stratego compiler `sc` by typing

```
> sc -i disjnf
```

Note that the suffix `.r` is not included in the name provided to the compiler. During compilation the compiler gives diagnostic output to indicate the progress of the compilation process.

```
> sc -i disjnf
```

```
...
```

```
compilation succeeded
```

Compilation creates the C program `disjnf.c` and the executable program `disjnf`.

USING COMPILED SPECIFICATIONS

The executable `disjnf` can be applied to a term to transform it. For example, the formula in file `prop.trm` is transformed by typing:

```
> disjnf < prop.trm
```

```
Or(Or(Atom("A"),Or(Not(Atom("B")),Not(Atom("C")))),Not(Atom("D")))
```

If rewriting fails the output would read

```
rewriting failed
```

```
module disjnf
imports lib prop-simplify
strategies

  main = stdio(eval; desugar; disj-nf)
```

Figure 23: Module defining `main` strategy.

DETECTING ERRORS IN STRATEGO SPECIFICATIONS

The errors made in Stratego programs can be classified into syntactic errors and logical errors. Logical errors are errors in the intended semantics of the specification. Syntactic errors are violations of the syntax or static semantics of the language. Not all syntactic errors are currently caught by the front-end of the compiler.

If compilation fails there is something wrong with the specification. The compiler should give an error message to indicate what is wrong. The following types of errors may occur.

UNDEFINED OPERATOR

If we forget the import of the library module `prop-truth`, we get the error message

```
error: operator T1 undefined
```

that indicates that the rule T1 is not defined.

UNBOUND VARIABLE

If in rule A9 we would make a typo:

```
A0 : Or(Or(x, y), z) -> Or(x, Or(y, s))
```

i.e., `s` instead of `z` in the right-hand side, we get the error message

```
error in rule A0 : variable s: used, but not bound
```

that indicates that the rule is trying to build a term with an uninstantiated variable.

COMMON ERRORS

The following are common errors

- strategy operator with wrong number of arguments
- term constructor with wrong number of arguments (run-time error)
- using term where strategy is expected (failure because of non-applicable congruence, undefined operator if some of the constructors are not defined in the signature)
- using a strategy where a term is expected

DEBUGGING STRATEGO PROGRAMS

In order to detect logical errors in specification it sometimes helps to understand what goes wrong at run-time (in addition to carefully reading the specification).

SYNTACTIC ERRORS WITH RUN-TIME MANIFESTATION

- the transformation crashes mentioning an assertion violation
- the transformation crashes with a segmentation fault

these are caused by shortcomings in the implementation

LOGICAL ERRORS

- the transformation fails while it is expected to succeed
- the transformation succeeds, but with a wrong result

SPLITTING UP THE SPECIFICATION

Often it is useful to test components of a specification to narrow down the place where the error occurs

- comment out tail end of a transformation pipeline
- narrow down the subject term

THE DEBUG OPERATOR

The `debug` operator prints the current term to `stderr`.

THE PRINTSTACK OPERATOR

The primitive `print-stack` prints the top `n` elements of the stack if applied to an integer term (as in `where(<print-stack> n)`) or the entire stack if applied to a non-integer term.

EXTENDING YOUR UNDERSTANDING OF STRATEGO

This tutorial only scratches the surface of Stratego specification. To extend your understanding of Stratego refer to the other existing documentation and bug me with requests to extend this tutorial.

OTHER DOCUMENTATION

Refer to *The Stratego Reference Manual* for an overview of all the language constructs.

Refer to *The Stratego Library* for a large number of standard strategies. The modules in the library are self documenting and often contain examples of the application of the strategies.

Refer to *The Stratego Compiler* for an overview of the implementation of Stratego and all the details of the compilation process.

PUBLICATIONS

There are a number of publications on the design, implementation and use of Stratego. See the bibliography for references. While these publications do not always reflect the current state of the language, they do reflect its spirit.

TOPICS FOR FUTURE INCLUSION IN THIS TUTORIAL

Here are some topics that I consider for inclusion in this tutorial. Which ones I tackle first will depend on popular demand and my mood.

- Discussion of (advanced) language constructs
- Primitives
- Architectures for program transformation
- Connecting Stratego components to other systems

BIBLIOGRAPHY

- [1] Mark G. J. Van den Brand and Eelco Visser. Generation of formatters for context-free languages. *ACM Transactions on Software Engineering and Methodology*, 5(1):1–41, January 1996.
- [2] Bas Luttik and Eelco Visser. Specification of rewriting strategies. In M. P. A. Sellink, editor, *2nd International Workshop on the Theory and Practice of Algebraic Specifications (ASF+SDF'97)*, Electronic Workshops in Computing, Berlin, November 1997. Springer-Verlag.
- [3] Eelco Visser. Strategic pattern matching. In P. Narendran and M. Rusinowitch, editors, *Rewriting Techniques and Applications (RTA'99)*, volume 1631 of *Lecture Notes in Computer Science*, pages 30–44, Trento, Italy, July 1999. Springer-Verlag.
- [4] Eelco Visser. *The Stratego Compiler*. Institute of Information and Computing Sciences, Universiteit Utrecht, Utrecht, The Netherlands, 1999.
- [5] Eelco Visser. *The Stratego Library*. Institute of Information and Computing Sciences, Universiteit Utrecht, Utrecht, The Netherlands, 1999.
- [6] Eelco Visser. *The Stratego Reference Manual*. Institute of Information and Computing Sciences, Universiteit Utrecht, Utrecht, The Netherlands, 1999.
- [7] Eelco Visser and Zine-el-Abidine Benaissa. A core language for rewriting. *Electronic Notes in Theoretical Computer Science*, 15, September 1998. In C. Kirchner and H. Kirchner, editors, *Proceedings of the Second International Workshop on Rewriting Logic and its Applications (WRLA'98)*, Pont-à-Mousson, France.
- [8] Eelco Visser, Zine-el-Abidine Benaissa, and Andrew Tolmach. Building program optimizers with rewriting strategies. *ACM SIGPLAN Notices*, 34(1):13–26, January 1999. *Proceedings of the International Conference on Functional Programming (ICFP'98)*.